

Should I work at a frontier lab?

Let's say you're talented and safety-conscious. What is the impact of working at a frontier lab? We can imagine a few scenarios, based on the proportion of safety-conscious employees:

Direct Impact

1 - Literally No One Is Safety-Conscious

This seems pretty bad. By default we expect this lab to just take the fastest path to AGI, aligned or not. No time for evals, safety training, or even mundane stuff such as not generating inappropriate images.

2 - Literally One Person Is Safety-Conscious

This seems kind of unstable. If this one person doesn't have an *extremely* strong will, the pressure to fit in and, like, not get fired will reduce us to Scenario 1. But if they do, it's a meaningful step up: now you have a whistleblower available to nip the worst outcomes in the bud.

3 - A Small Group (<5%) Of People Are Safety-Conscious

This is better than 2, because with friends, the safety-conscious people have more support, just psychologically. But this doesn't change the fact that these people will have a very hard time getting promoted, or even avoiding being fired.

4 - A Significant Minority Of People Are Safety-Conscious

Now we're getting to the point where the safety-conscious people can actually *do* something. Without the safety-conscious people, the organization actually takes enough of a hit that they can't just fire everyone they don't like. I expect this many safety-conscious people to, for example, be able to make some noise about Grok generating explicit images on demand, and maybe implement cheap safeguards, but not necessarily take it down.

5 - Around Half Of People Are Safety-Conscious

Now we have enough safety-conscious people that some of them are probably in positions of power. The more safety-conscious people, proportionally the more of them will take power,

because they will have a good coalition. If your positions are more popular, you're more likely to get appointed. At this point we get the potential for taking down Grok's explicit-image-making capabilities, but not necessarily more drastic actions such as saying no to the Department of War.

The specific percentage matters too. At 25%, I'd expect stuff like "don't train on the CoT" to be in the Overton window but difficult. At 75%, it's clear as day, and stuff like "saying no to the DoW" becomes conscionable, even if hard.

6 - Most People Are Safety-Conscious

I actually don't think there's too much of a difference between "most" and "all", because at this point the safety-conscious people have a coalition and can do whatever they want. Both "most" and "all" have the problem of guarding against potential future bad actors ("most" also has current bad actors, so maybe a little different). At this point we get stuff like saying no to the DoW and delaying frontier model releases. That being said, even at 100% it's not a free lunch: the firm still faces race dynamics and market/investor pressures. It's maybe even plausible that due to race dynamics, even at a high percentage your work is still net-negative. I feel like at close to 100% though, the people are dedicated enough to figure something out.

- 0% - full race
- 0-5% - full race, but with whistleblowing
- 5-20% - full race, but don't do anything too obviously stupid
- 20-80% - gradually taking more safe actions
- 80-100% - do the best we can modulo market pressures and race dynamics

I imagine a big jump from 0% to 5% safety-conscious people, and fairly linear from there.

Seat At The Table

Another factor to consider is that, by joining an AI org, you're changing the percentage yourself! For example, if an AI org is 40% safety-conscious, and you join it, then assuming you're safety-conscious, you've counterfactually increased the number of safety-conscious employees by 0.6 in expectation, since your replacement would only be 40% likely to be safety-conscious.

More importantly, if you do a good job and up having influence over the company's decisions, this can be a huge force for good. So if you're a 10x engineer, or very charismatic, you could change the direction of the company. This is mainly a factor later on in your career once you take on higher positions, but can still be important early on, especially if you're ambitious and do good work.

Value Drift

I hinted at this in the section about only 1 safety-conscious person, but you might not end up the same person you started as. I can easily imagine myself being brainwashed into thinking "oh but the best thing to do for safety is to work on capabilities, to prevent a capabilities overhang!". This also impacts any future AI-safety related work you do (or would have done). This is most likely in precisely the scenarios where you're needed most. So we can roughly say your direct impact is:

- At 0-5%: huge if you have a very strong will, otherwise negative
- At 5-20%: your direct work might be negative, but you push the firm to a more positive number (assuming your will is strong)
- At 20-80%: same, but your direct work is more likely to be positive
- At 80-100%: you're adding force to an already positive movement, as long as you keep race dynamics in consideration.

Your early-career years have the most future years ahead of you for your values to influence your impact, and your values tend to stabilize as you get older, making this a larger consideration earlier on (though not negligible later though).

Career Capital

If you're early on in your career, your direct work now isn't going to be that impactful (compared to more experienced people now, and yourself later). The main effect of your work will actually be accumulating career capital: transferable skills, connections, and something to put on your resume. For example, if you work at Anthropic, you're probably not developing the breakthrough to achieve AGI, but you'll get a lot closer to doing so, as well as having much more legible achievements to some other AI organization you might want to work for in the future, and the connections to find these opportunities.

This means, early on, the main considerations are career capital accumulation vs value drift, as opposed to your direct work or having a seat at the table, which are much more important if you're senior. So, assuming relatively long timelines, you probably want to work at the company that maximizes (career capital force multiplier) × (alignment to your current values in the future).

Plausibly, the top frontier labs (e.g. OpenAI, Anthropic) give you huge career capital force multipliers. If you worked there, basically anyone will want to hire you, and you'll be damn good at your job. So, you just need to make sure the second term is positive. That is, if your values remain intact, you'll use the career capital you built at your first job to do great things with AI in a few years. But if not, you could even end up making a net-negative impact that your career capital multiplies, arguably the worst possible combination. Maintaining your values in the face of social pressure is harder than it looks: it's easy to think you have a strong will, but rationalize your worries away once the incentives stack against you. Just look what happened to Sam

Altman! He started out aggressively pro-AI safety and regulation, and now he's lobbying for state bills giving [OpenAI immunity against liability for catastrophic harms](#). Why would you want immunity against liability for catastrophic harms? Hmm.....

Also, I've experienced value drift in my own life. Two years ago, I swore to myself that I'd never join the dark side and work at a quant finance company. Now, I have an offer at Jane Street. Of course, right now I've rationalized it by saying I want to build career capital. Maybe this is true! But I'm still wary of getting value-drifted further.

But What About The Counterfactual?

This is kinda tricky. On the one hand, everyone wants a position at these frontier AI labs, so you might think this would make you very replaceable. But also, these labs are ravenous for talent. The synthesis is that finding *hires* is easy, but finding *good hires* is hard. We should think about both your delta from the next hire taken, along with taking into account that without you they might not have sized up.

- At 0-20%: your replacement probably doesn't have that strong of a will, even if safety-conscious. So if you do, huge counterfactual impact right there. You'd be preventing someone else from getting value drifted.
- At 20-80%: this dilutes the "direct work negative" argument (or positive, if closer to the 80 side), while keeping the "seat at the table" argument. So, assuming that you can stay aligned, you should really be thinking about the seat at the table that you're replacing. Exception: if you're a 10x engineer late in your career, you should also be thinking about your direct work.
- At 80-100%: you want to make sure you're either a 10x engineer, or that the org you're working for actually can use more people. Otherwise, you're replacing someone just as aligned and capable as you are.

Counterfactual impact throws a wrench into career capital considerations as well: if you're early-career and being hired into an early-career role, you don't just need to make use of the connections, credentials, and skills you've gained. You need to make better use of them than the next person in line, assuming both of you would make a positive impact with those advantages later on. However, if you think the person you're replacing is less likely to do so, or if they're more likely to get value drifted, your counterfactual impact remains positive.

Summary

Scenario	Early Career	Experienced
Safety-conscious organization (e.g. Redwood, Anthropic?)	Join to build career capital, but must be ambitious (since the marginal hire also would've used it well)	Join if you're a 10x engineer and do great work, just be wary of race dynamics
Somewhere in the middle (e.g. Deepmind, OpenAI?)	Join to build career capital, but try not to get value drifted	Join to take a seat at the table
Non-safety-conscious organization (e.g. xAI?)	Steer clear unless very strong-willed and lack better options to build career capital	Join to take a seat at the table if strong-willed, and potentially whistleblow

References

[Reasons for and against working on technical AI safety at a frontier AI lab — EA Forum](#)

[Aspiring AI Safety Researchers: Consider “Atypical Jobs” in the Field Instead. — EA Forum](#)

[Why I think it's net harmful to do technical safety research at AGI labs — LessWrong](#)

[Anonymous advice: If you want to reduce AI risk, should you take roles that advance AI capabilities? | 80,000 Hours](#)

[Making commitments to yourself to avoid System 1 traps](#)

[Should you work at a leading AI lab? \(including in non-safety roles\) — LessWrong](#)

[Shubs \(contributed seat at the table argument\)](#)