Proposal for the National AI Security Organization (NASO)

The National AI Security Organization incorporates independent AI safety researchers, biosecurity experts, cybersecurity experts, economists, and government appointed representatives to establish pre-deployment security criteria that AI labs have to follow before releasing new models to the public. New models will be red teamed against these criteria by a group of rival lab safety researchers with API level access before the model is released to the public. The red team has the ability to escalate the testing to a more invasive evaluation if they find that the model meets criteria established by the independent board for capability thresholds. The model is then evaluated more thoroughly by the independent board with full access to the model.

Background

The current AI policy landscape is criticized for being fragmented. With the main pieces of legislation being the RAISE Act in New York[1] and SB-53 in California.[2] These pieces of legislation are criticized on many dimensions including regulatory capture, compliance burdens, and being ineffective.[3] The federal response with preemption challenges the fragmentation but does not provide substantive solutions to address AI risk.[4]

The UK AI Safety Institute—the largest government evaluation program—has admitted that meaningful evaluation requires deeper access than labs have been willing to provide, including access to models with safety filters disabled, fine-tuning API access, and internal evaluation results.[5] The race dynamics between labs have meant that labs have gotten less strict in their evaluations rather than more strict—Anthropic dropped its commitment to halt development if safety measures couldn't be guaranteed in its February 2026 RSP revision, explicitly citing competitive pressure from rivals who are "blazing ahead."[6] Therefore, it is critical that the incentives of the frontier labs be changed so that mitigating risk is not a liability but an asset.

Implementation

Organizational Structure: The National AI Security Organization board will be composed of thirteen members: five representatives from AI labs—four from the top AI labs as determined by revenue and a fifth member chosen from another lab that the remaining members of the board deem noteworthy but haven't commercialized, five independent AI safety researchers appointed by the National Academies of Sciences, Engineering, and Medicine (NASEM),[7] three government appointees, including representatives from Commerce, DHS, and the DOD. The role of the DOC appointee is to provide perspective on the Economic implications of AI advancements, anchored in NIST's existing AI Risk Management Framework.[8] The member representing the DOD will be responsible for the national security perspective. The representative for the DHS will provide insight on AI advancements and domestic threats. The AI lab representatives will be reappointed each year, the AI safety researchers will be appointed every two years, and the remaining three members are subject to change depending on the current president's appointment and approval by Congress.

The Risk Standards Board—composed of AI safety researchers, held to the same standards as the governing board, biosecurity and cybersecurity experts, and economists—will work under the governing board and will be appointed by the governing board with a simple majority approval. They will be responsible for creating an evolving standard—rewritten every quarter—with which red teams will evaluate new models. The standard will include categories such as automated research, cyber-offense, bioweapon capabilities, economic implications, and deceptive alignment. Specifically the standard will include a framework for weighing risk against deployment value, as some risk is worth taking given the right upside. This approach draws on Jones (2023), which models the optimal deployment of AI given the tradeoff between unprecedented economic growth and existential risk, demonstrating that a social planner's tolerance for risk depends on the magnitude of growth benefits, the level of threat, and the curvature of the utility function.[9] These standards will be publicly available for transparency and more democratic accountability, following the notice-and-comment model used in federal administrative rulemaking.

The red team will consist of researchers from AI labs. When an AI company above $100 million in annual revenue from frontier model development or deployment wants to release a new model they will be assigned a red team to test their model in an air gapped environment for a period of two weeks. The governing board also has the ability to add to the program any lab that they determine to be on par with the frontier labs.The red team will have API level access to the new model so they can test the new model on the criteria defined by the Risk Standards Board. The red team then presents their findings to the governing board for approval. If they notice unexpected capability increases they will escalate to a more invasive evaluation. An **escalated evaluation** will be carried out by the Risk Standards Board to prevent IP theft.[11] The board can then present their findings to the governing board, who can either approve the model for public release or issue a mandatory market access ban—the model may not be deployed, licensed, or made accessible via API within the United States until deficiencies are remediated and the model passes re-evaluation. For models receiving a critical risk finding, the NASO refers findings to the Bureau of Industry and Security for export control determination and notifies the Executive and Legislative branch for further actions like a pause on development and negotiations for an international pause. Deploying a model without NASO approval is considered a civil violation with penalties proportional to the lab's revenue.

Tradeoffs

Red teaming requirements will slow down innovation by diverting lab resources and slowing the speed of deployment, but the slowdown is imposed on all major domestic labs. The requirements could harm competitiveness with labs in other countries, however the standards can be adjusted to reflect competitive dynamics as needed.[12] There is also some risk that frontier models that are deemed unexpectedly capable might be leaked if they are exposed to compromised individuals in the escalated evaluation. It might be difficult to convince talented AI safety researchers to work for the NASO given the six-year employment restriction. AI labs being included in the governing board introduces risks of regulatory capture, but it is important for the viability of the policy, ensuring sufficient expertise and preventing political pushback. The speed at which the evaluation criteria made by the Risk Standards Board is changed might be too slow if capabilities grow faster, but it is critical for criteria to be somewhat consistent so labs can work off them.

References

[1] New York RAISE Act, S6953B/A6453B (Dec. 2025). NYSenate.gov.

[2] California SB-53, Transparency in Frontier Artificial Intelligence Act (Sept. 2025). [3]

Brookings Institution, "What is California's AI safety law?" (Dec. 2025). brookings.edu.

[4] Executive Order, "Ensuring a National Policy Framework for Artificial Intelligence" (Dec. 11, 2025). WhiteHouse.gov.

[5] UK AI Safety Institute, "Early Lessons from Evaluating Frontier AI Systems" (Oct. 2024). aisi.gov.uk.

[6] "Exclusive: Anthropic Drops Flagship Safety Pledge," TIME (Feb. 2026); Anthropic, Responsible Scaling Policy v3.0 (Feb. 2026). anthropic.com/responsible-scaling-policy.

[7] The National Academies of Sciences, Engineering, and Medicine operate under an 1863 congressional charter and maintain a dedicated AI topic area, including a 2024 consensus report on AI and the workforce requested by Congress. nationalacademies.org.

[8] NIST AI Risk Management Framework (AI RMF 1.0, Jan. 2023); Generative AI Profile (July 2024). nist.gov.

[9] Jones, Charles I. "The A.I. Dilemma: Growth versus Existential Risk." NBER Working Paper No. 31837 (2023); published in American Economic Review: Insights, vol. 6(4), 2024.

[10] Epoch AI, "The combined revenues of leading AI companies grew by over 9x in 2023-2024" (April 2025) — no AI company beyond OpenAI, Anthropic, and Google DeepMind surpassed $100M in revenue in 2024 from selling access to their own models. epoch.ai.

[11] Longpre et al., "A Safe Harbor for AI Evaluation and Red Teaming" (2024). Knight First Amendment Institute, Columbia University.

[12] EU AI Act, Regulation (EU) 2024/1689; GPAI Code of Practice (2025) — demonstrates that proportional compliance requirements can coexist with a competitive industry.

[13] AI Lab Watch, "AI Companies Aren't Really Using External Evaluators" (May 2024). ailabwatch.org — documents the scarcity of qualified frontier AI evaluators.