

Policy Brief: Regulating Frontier AI Through A Modular Risk-Based Approach

Policy Recommendation: The United States should implement a federal Frontier AI Developer Accountability Act (“FAIDA”) that regulates *AI Entities* developing and deploying Frontier AI models, designated by a multi-factor threshold. This represents an innovation-based modular approach that is future-looking and preserves statutory and regulatory enforcement through a baseline AI safety framework.

Strategic Context: Frontier AI models (“FAI”) are versatile, general-purpose technologies, and have shown emergent unpredictable behavior. OpenAI’s o1 attempted to deactivate its oversight mechanism in 5% of test cases; Claude 3 Opus strategically faked alignment in 12%; DeepSeek R1 attempted to hack environments to win chess games.ⁱ Several AI Entities have flagged in their system cards that models are approaching critical capability for CBRN weaponsⁱⁱ, cybersecurity, and autonomous operation.ⁱⁱⁱ The risks arising from FAI can be put into two buckets: (i) risks with evidence of harm: scams, NCII, CSAM, bias, reliability issues, and privacy violations and (ii) risks with unclear yet growing evidence: CBRN capability, mass labor market impact, autonomous cybersecurity attacks, loss of developer control, and other emergent harms.

Regulation should be based on a dual-threshold coverage criteria that addresses the most significant risks through a tiered, modular approach. Overseen by an Advisory Committee and avoiding federal preemption, our proposal highlights baseline AI safety readiness compliance for the AI Entities. We build on the White House National Policy Framework for AI^{iv} and EOs^v issued on AI policy improving AI-readiness and incentivizing innovation.

The Frontier AI Developer Accountability Act (FAIDA):

FAIDA regulates the largest AI Entities building the most advanced AI systems. Instead of focusing on static thresholds based on compute (NY’s RAISE Act^{vi}, California’s SB-53^{vii}), or use-cases (EU AI Act^{viii}), the U.S. should regulate FAI across its supply chain through a regulatory ladder that scales with risk. FAIDA sets out the following regulatory framework: -

- I. **Entity-based Regulation:** US-based AI Entities developing and deploying FAI.
- II. **Coverage:** AI Entities meeting *at least two* of the following statutory thresholds would be covered: (i) Assessment of marginal risk through specified benchmarks; (ii) Compute capabilities at 10^{25} FLOPs; (iii) Annual AI-related R&D exceeds \$1 billion; or (iv) Annual AI-related Revenue exceeds \$500 million.
- III. **Regulatory Ladder:** Tiered approach to AI safety based on emergence of *marginal risk* from FAI. Tier 1 applies to all covered AI Entities. Tier 2 is triggered when an AI Entity’s Tier 1 assessment signals capabilities for: CBRN, automated cyber operations, autonomous behavior/self-replication, evasion of developer controls, or behavior that credibly threatens mass casualties or harm.

FOR: Senator Brian Schatz (D-HI), Senior Member, U.S. Senate Committee on Commerce, Science, and Transportation.

- IV. **Covered Entity Obligations: Tier 1 - Baseline Safety Framework (BSF)**; Publish Secure Development Frameworks (SDFs), system cards, safety incident reporting, risk assessment thresholds, summaries of catastrophic risk assessments, and demonstrate compliance with such disclosures and whistleblower protections. **Tier 2 - Heightened Oversight Framework (HOF)**; Demonstrate compliance with BSF, conduct ongoing risk monitoring, external red-teaming, heightened model weights security, and rapidly restrict model access in response to critical safety incidents.^{ix}
- V. **Preemption and Sunset Clause**: FAIDA recognizes that FAI *development* is inherently interstate and creates a clear unified standard for Tier 1 and 2 activities. States retain full authority to enforce *deployment*-based AI safety laws for consumer protection, children's safety, domain/sector-specific, etc. After five years, federal preemption ceases for Congressional reassessment.
- VI. **Open-weight models**: Before a covered AI Entity releases open-weight models, it must show compliance with BSF requirements (system card evals, deployment guardrails). If pre-release evaluations trigger Tier 2 capabilities, it must take active measures to irreversibly degrade identified capabilities^x (for e.g., CBRN).
- VII. **Rebuttable Safe Harbor**: If a covered AI Entity demonstrates implementation of BSF and HOF, it receives safe harbor for third-party downstream harms.
- VIII. **Enforcement Authority**: FTC would enforce misrepresentation in BSF and HOF disclosures. State AGs and statutorily-authorized bodies to continue enforcement as per state laws.
- IX. **Establishment of Frontier AI Advisory Committee (FAAC)^{xi} and Warranting NIST AI Safety Institute's Standards**: Appointed by the Secretary of Commerce and comprising frontier AI researchers, third-party evaluators, civil society representatives, and at least one national security expert, with no more than one-third drawn from covered AI Entities, advises the Secretary and informs NIST standards applicable to AI Entities. Mandate: (i) recommend revisions to FAIDA's coverage criteria and marginal risk benchmarks as capabilities evolve; (ii) review BSF and HOF disclosures and flag systemic gaps or underreporting to the FTC and Secretary of Commerce; and (iii) publish an 'Annual Frontier AI Risk Report'. The NIST AI Safety Institute will develop and update voluntary technical benchmarks for assessing marginal risk and SDF guidance.
- X. **Severability**: If one provision is invalidated, other sections of the statute remain valid.

Trade-offs:

FAIDA places limits on harms arising from agentic systems. However, autonomous multi-step action without human involvement may be addressed through complementary legislation. It is acknowledged that statutorily defined Tier 2 critical harms risk obsolescence as capabilities advance. This tradeoff represents the inherent tension in incentivizing innovation that also carries enormous risks. FAIDA seeks to address these risks through its statutory levers and obligations.

FOR: Senator Brian Schatz (D-HI), Senior Member, U.S. Senate Committee on Commerce, Science, and Transportation.

ⁱ Rishi Bommasani et al., *The California Report on Frontier AI Policy* (Joint California Policy Working Group on AI Frontier Models, June 17, 2025), <https://www.gov.ca.gov/wp-content/uploads/2025/06/June-17-2025---The-California-Report-on-Frontier-AI-Policy.pdf>.

ⁱⁱ CBRN stands for Chemical, Biological, Radiological, and Nuclear, a classification used to describe weapons capable of causing mass casualties and widespread destruction.

ⁱⁱⁱ Rishi Bommasani et al., *The California Report on Frontier AI Policy* (Joint California Policy Working Group on AI Frontier Models, June 17, 2025), <https://www.gov.ca.gov/wp-content/uploads/2025/06/June-17-2025---The-California-Report-on-Frontier-AI-Policy.pdf>:

OpenAI assessed risks in areas like CBRN, cybersecurity, and autonomous behavior as “medium”. Anthropic warned that its next models might require stringent safeguards (ASL-3), particularly due to their potential to assist even moderately skilled individuals in developing or deploying CBRN weapons and later implemented such safeguards for Claude Opus 4 out of caution. Similarly, Google indicated that its Gemini 2.5 Pro had already crossed early warning thresholds, suggesting that near-future versions could reach “critical capability” levels, especially in enabling highly impactful cyberattacks with significantly reduced resources.

^{iv} The White House, *National Policy Framework for Artificial Intelligence: Legislative Recommendations* (Washington, DC: The White House, March 2026), <https://www.whitehouse.gov/wp-content/uploads/2026/03/03.20.26-National-Policy-Framework-for-Artificial-Intelligence-Legislative-Recommendations.pdf>.

^v Executive Order 14110, 88 Fed. Reg. 75191 (November 1, 2023), <https://www.federalregister.gov/documents/2023/11/01/2023-24283/safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence> & The White House, *Executive Order 14179: Removing Barriers to American Leadership in Artificial Intelligence* (January 23, 2025), <https://www.whitehouse.gov/presidential-actions/2025/01/removing-barriers-to-american-leadership-in-artificial-intelligence/>.

^{vi} New York State Assembly, *Responsible AI Safety and Education Act (RAISE Act)*, A.6453-A, 2025–2026 Reg. Sess. (introduced March 5, 2025), <https://www.nysenate.gov/legislation/bills/2025/A6453/amendment/A>.

^{vii} California Legislature, *Transparency in Frontier Artificial Intelligence Act*, Senate Bill 53, 2025–2026 Regular Session, approved by Governor September 29, 2025, Chapter 138, Statutes of 2025, https://leginfo.legislature.ca.gov/faces/billTextClient.xhtml?bill_id=202520260SB53.

^{viii} European Parliament and Council of the European Union, *Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act)*, Official Journal of the European Union L 2024/1689 (July 12, 2024), <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng>.

^{ix} Anita Rao et al., *Challenges to the Monitoring of Deployed AI Systems*, NIST AI 800-4 (Gaithersburg, MD: National Institute of Standards and Technology, March 2026), <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.800-4.pdf>: NIST’s most recent paper points to the limitations of pre-deployment evaluations to not fully capture emergent model behavior in real-

FOR: Senator Brian Schatz (D-HI), Senior Member, U.S. Senate Committee on Commerce, Science, and Transportation.

world scenarios. Instead, it emphasizes the growing importance of post-deployment monitoring and introduces a taxonomy of six monitoring categories.

^x Measures include validated unlearning and tamper-resistance techniques, *See* Fan Liu et al., *Rethinking Machine Unlearning for Large Language Models*, 7 *Nature Mach. Intelligence* 208 (2025), <https://doi.org/10.1038/s42256-025-00985-0>; Nathaniel Li et al., *The WMDP Benchmark: Measuring and Reducing Malicious Use With Unlearning*, 41st International Conference on Machine Learning (ICML 2024) (2024), <https://arxiv.org/abs/2403.03218>; Rishub Tamirisa et al., *Tamper-Resistant Safeguards for Open-Weight LLMs*, 13th International Conference on Learning Representations (ICLR 2025) (2024). However, durably safeguarding open weight models against CBRN capabilities is an active area of research and is still evolving, *See* A. Feder Cooper et al., *Machine Unlearning Doesn't Do What You Think: Lessons for Generative AI Policy and Research*, 39th Conference on Neural Information Processing Systems (NeurIPS 2025) (2024), <https://arxiv.org/abs/2412.06966>.

^{xi} Subject to the Federal Advisory Committee Act (FACA; 5 U.S.C. Chapter 10): Meghan M. Stuessy and Kathleen E. Marchsteiner, *The Federal Advisory Committee Act (FACA): Overview and Considerations for Congress*, Congressional Research Service, R47984 (Washington, DC: Congressional Research Service, 2024), <https://www.congress.gov/crs-product/R47984>.