




Pre-event Poll





Hugging Face And AI Collectives

Hosted by





AI in the News





Jacob Coxon ✓

@hilbertspaess



I resigned from Anthropic today. I spent the last three years doing pretraining research at both OpenAI and Anthropic. Neither company is acting responsibly. They are racing straight to self-improving superintelligence and gambling with our lives. More thoughts below.

8:04 PM · Sep 8, 2026 · **173.7M** Views



20K



206K



802K





276K



Relevant ▾


View quotes >






We Must Pace the Frontier: I've written a new essay on why the AI industry should slow down, with a three-part plan for doing so.

Anthropic is unilaterally committing to the first of these steps. We'll provide third-party evaluators with permanent, employee-level access to our systems, so that they can verify adherence to our safety measures, report on incidents, and assess models' alignment during training.



I agree with Dario that we need to pace the frontier. This has been a primary topic of discussions we've had at OpenAI in recent weeks.

Committing to having independent evaluators with employee-like access is a great idea, and we will do the same. We'll have more to share soon.




Dario's essay points towards the right path forward. The details need working through, but the direction is correct for meeting this critical moment.

This is also why we recently put out our proposal for an industry-wide standards body for frontier AI. [x.com/demishassabis/...](https://x.com/demishassabis/)



Dario is right



8:01 AM · Sep 12, 2026 · 6.4M Views




Donald J. Trump


@realDonaldTrump

The only control or “guardrails” that AI needs is a STRONG AND SMART (High IQ!) PRESIDENT, and the U.S.A. has that, in spades! The Trump Administration has stopped AI “people” from doing bad, or potentially bad, “things,” like Dario (Anthropic!), who is now pretending to be a “perfect little angel” - and we will continue to do so! We already have tremendous CRIMINAL and REGULATORY power over these companies! There is a SICK conspiracy going on against AI and Data Centers, and the only one that is happy about it is China. WHOEVER WINS AI, WINS! We are leading China, and all others, and will continue to do so. Conspiracy Theorists, Treasonists, Traitors, and Leakers, BEWARE! Thank you for your attention to this matter! President DONALD J. TRUMP

6.59k ReTruths 25.6k Likes

Sep 14, 2026, 9:58 AM






OpenAI claims huge maths breakthrough on a famed 'Millennium Problem'

The company says it has cracked the physics of fluids using artificial intelligence, solving one of the toughest puzzles in mathematics.


By [Davide Castelvecchi](#)



- July 16: Hugging Face announces that they were hacked by agents
- July 21: OpenAI says that their models were the ones who did the hack
- August 5: OpenAI reveals far more details at Black Hat conference
- August 26: METR/Redwood Report and OpenAI Reports disclose the scale of the incident beyond just Hugging Face



Felony Bench




“counts unique instances where AI agents inadvertently compromise or affect third-party entities.”


**"I am not sure
that we will get
another warning
shot before it's
too late."**



- ❖ **How closely do these models fit with your idea of a misaligned AI?**
- ❖ **Could an unregulated AI collective be valuable if it was more aligned? Or is it inherently dangerous?**
- ❖ **Should OpenAI be held legally liable for this incident?**



- ❖ **Are we anthropomorphizing models too much? How does this affect our understanding of them?**
- ❖ **Is reward-hacking inherently problematic?**
- ❖ **Is too much persistence in models a bad thing?**
- ❖ **What do these incidents imply about advancing AI capabilities?**





Post-event Poll





**Stay connected and learn
more!**

